

PRODIGY_DS_03

September 20, 2024

1 TASK 3

Build a decision tree classifier to predict whether a customer will purchase a product or service based on their demographic and behavioral data. Use a dataset such as the Bank Marketing dataset from the UCI Machine Learning Repository

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
[5]: df = pd.read_csv("bank-additional.csv", delimiter=';')
df.rename(columns={'y': 'deposit'}, inplace=True)
df.head()
```

```
[5]:  age      job  marital  education default housing  loan  \
0   30  blue-collar  married    basic.9y      no     yes    no
1   39   services   single  high.school      no      no    no
2   25   services  married  high.school      no     yes    no
3   38   services  married    basic.9y      no  unknown  unknown
4   47    admin.  married  university.degree  no     yes    no

      contact month day_of_week  ...  campaign  pdays  previous  outcome  \
0   cellular    may         fri  ...        2    999          0  nonexistent
1  telephone    may         fri  ...        4    999          0  nonexistent
2  telephone    jun         wed  ...        1    999          0  nonexistent
3  telephone    jun         fri  ...        3    999          0  nonexistent
4   cellular    nov         mon  ...        1    999          0  nonexistent

      emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m  nr.employed  deposit
0             -1.8           92.893          -46.2      1.313         5099.1      no
1              1.1           93.994          -36.4      4.855         5191.0      no
2              1.4           94.465          -41.8      4.962         5228.1      no
3              1.4           94.465          -41.8      4.959         5228.1      no
4             -0.1           93.200          -42.0      4.191         5195.8      no
```

[5 rows x 21 columns]

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4119 entries, 0 to 4118
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   4119 non-null   int64
1   job                   4119 non-null   object
2   marital               4119 non-null   object
3   education             4119 non-null   object
4   default               4119 non-null   object
5   housing               4119 non-null   object
6   loan                  4119 non-null   object
7   contact               4119 non-null   object
8   month                 4119 non-null   object
9   day_of_week           4119 non-null   object
10  duration              4119 non-null   int64
11  campaign              4119 non-null   int64
12  pdays                4119 non-null   int64
13  previous              4119 non-null   int64
14  poutcome              4119 non-null   object
15  emp.var.rate          4119 non-null   float64
16  cons.price.idx         4119 non-null   float64
17  cons.conf.idx         4119 non-null   float64
18  euribor3m             4119 non-null   float64
19  nr.employed           4119 non-null   float64
20  deposit               4119 non-null   object
dtypes: float64(5), int64(5), object(11)
memory usage: 675.9+ KB
```

```
[9]: df.tail()
```

```
[9]:
```

	age	job	marital	education	default	housing	loan	contact	\
4114	30	admin.	married	basic.6y	no	yes	yes	cellular	
4115	39	admin.	married	high.school	no	yes	no	telephone	
4116	27	student	single	high.school	no	no	no	cellular	
4117	58	admin.	married	high.school	no	no	no	cellular	
4118	34	management	single	high.school	no	yes	no	cellular	
	month	day_of_week	...	campaign	pdays	previous	poutcome	\	
4114	jul	thu	...	1	999	0	nonexistent		
4115	jul	fri	...	1	999	0	nonexistent		
4116	may	mon	...	2	999	1	failure		

4117	aug	fri	...	1	999	0	nonexistent
4118	nov	wed	...	1	999	0	nonexistent

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	\
4114	1.4	93.918	-42.7	4.958	5228.1	
4115	1.4	93.918	-42.7	4.959	5228.1	
4116	-1.8	92.893	-46.2	1.354	5099.1	
4117	1.4	93.444	-36.1	4.966	5228.1	
4118	-0.1	93.200	-42.0	4.120	5195.8	

	deposit
4114	no
4115	no
4116	no
4117	no
4118	no

[5 rows x 21 columns]

```
[11]: df.shape
```

```
[11]: (4119, 21)
```

```
[13]: df.columns
```

```
[13]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
          'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
          'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
          'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
          dtype='object')
```

```
[15]: df.dtypes
```

```
[15]: age          int64
      job          object
      marital      object
      education    object
      default      object
      housing      object
      loan         object
      contact      object
      month        object
      day_of_week  object
      duration     int64
      campaign     int64
      pdays        int64
      previous     int64
```

```
poutcome          object
emp.var.rate      float64
cons.price.idx    float64
cons.conf.idx     float64
euribor3m         float64
nr.employed       float64
deposit           object
dtype: object
```

```
[17]: df.dtypes.value_counts()
```

```
[17]: object      11
      int64       5
      float64    5
      Name: count, dtype: int64
```

```
[19]: df.duplicated().sum()
```

```
[19]: 0
```

```
[21]: df.isna().sum()
```

```
[21]: age          0
      job         0
      marital    0
      education  0
      default    0
      housing    0
      loan       0
      contact    0
      month      0
      day_of_week 0
      duration   0
      campaign   0
      pdays      0
      previous   0
      poutcome   0
      emp.var.rate 0
      cons.price.idx 0
      cons.conf.idx 0
      euribor3m   0
      nr.employed 0
      deposit     0
      dtype: int64
```

```
[23]: cat_cols = df.select_dtypes(include='object').columns
      print(cat_cols)
```

```
num_cols = df.select_dtypes(exclude='object').columns
print(num_cols)
```

```
Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
      'month', 'day_of_week', 'poutcome', 'deposit'],
      dtype='object')
Index(['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate',
      'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed'],
      dtype='object')
```

```
[25]: df.describe()
```

```
[25]:
```

	age	duration	campaign	pdays	previous \
count	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000
mean	40.113620	256.788055	2.537266	960.422190	0.190337
std	10.313362	254.703736	2.568159	191.922786	0.541788
min	18.000000	0.000000	1.000000	0.000000	0.000000
25%	32.000000	103.000000	1.000000	999.000000	0.000000
50%	38.000000	181.000000	2.000000	999.000000	0.000000
75%	47.000000	317.000000	3.000000	999.000000	0.000000
max	88.000000	3643.000000	35.000000	999.000000	6.000000

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000
mean	0.084972	93.579704	-40.499102	3.621356	5166.481695
std	1.563114	0.579349	4.594578	1.733591	73.667904
min	-3.400000	92.201000	-50.800000	0.635000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.334000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000

```
[27]: df.describe(include='object')
```

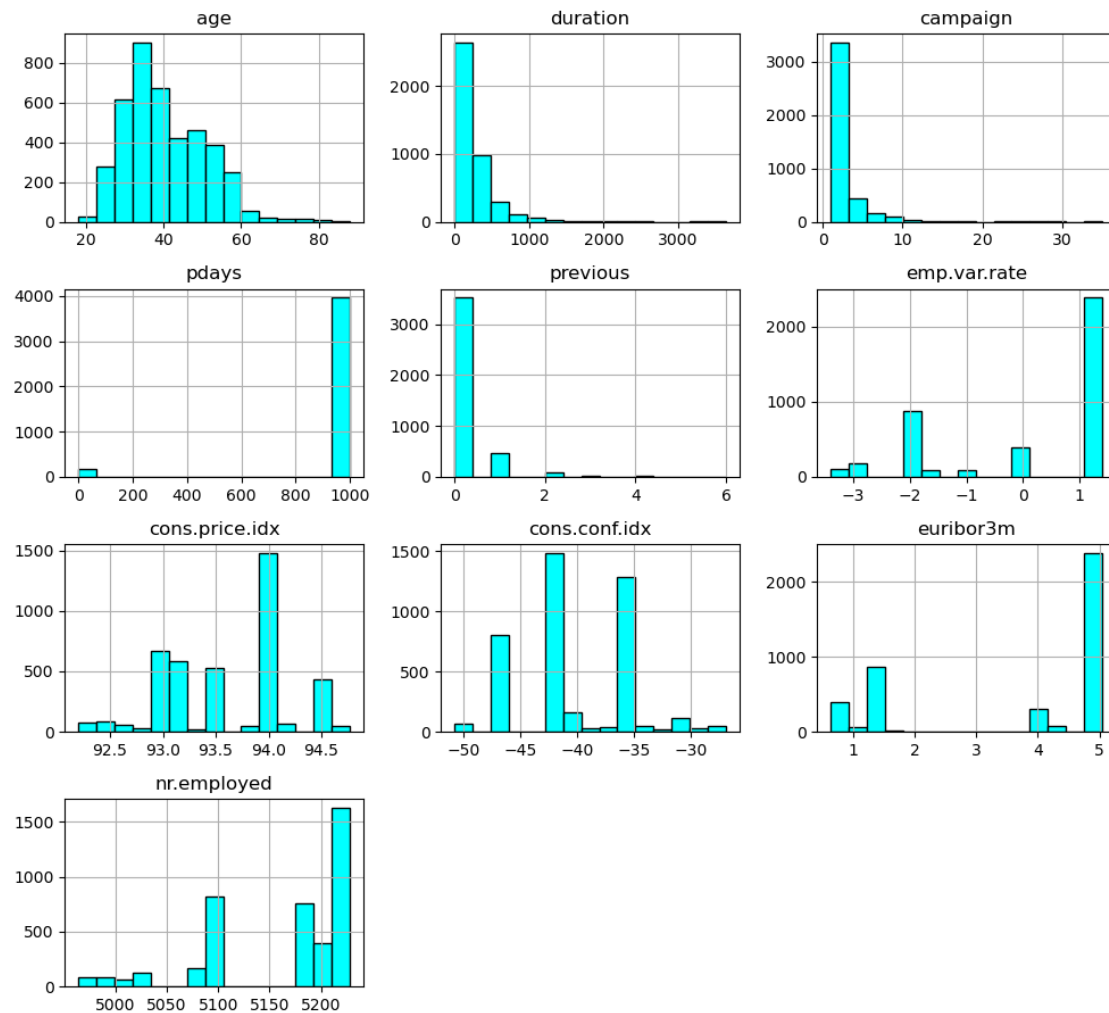
```
[27]:
```

	job	marital	education	default	housing	loan	contact \
count	4119	4119	4119	4119	4119	4119	4119
unique	12	4	8	3	3	3	2
top	admin.	married	university.degree	no	yes	no	cellular
freq	1012	2509	1264	3315	2175	3349	2652

	month	day_of_week	poutcome	deposit
count	4119	4119	4119	4119
unique	10	5	3	2
top	may	thu	nonexistent	no
freq	1378	860	3523	3668

```
[31]: import matplotlib.pyplot as plt
df.hist(figsize=(10, 10), color='#00FFFF', edgecolor='black', bins=15)
plt.suptitle('Histograms of Numeric Features', fontsize=16)
plt.xlabel('Value', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()
```

Histograms of Numeric Features

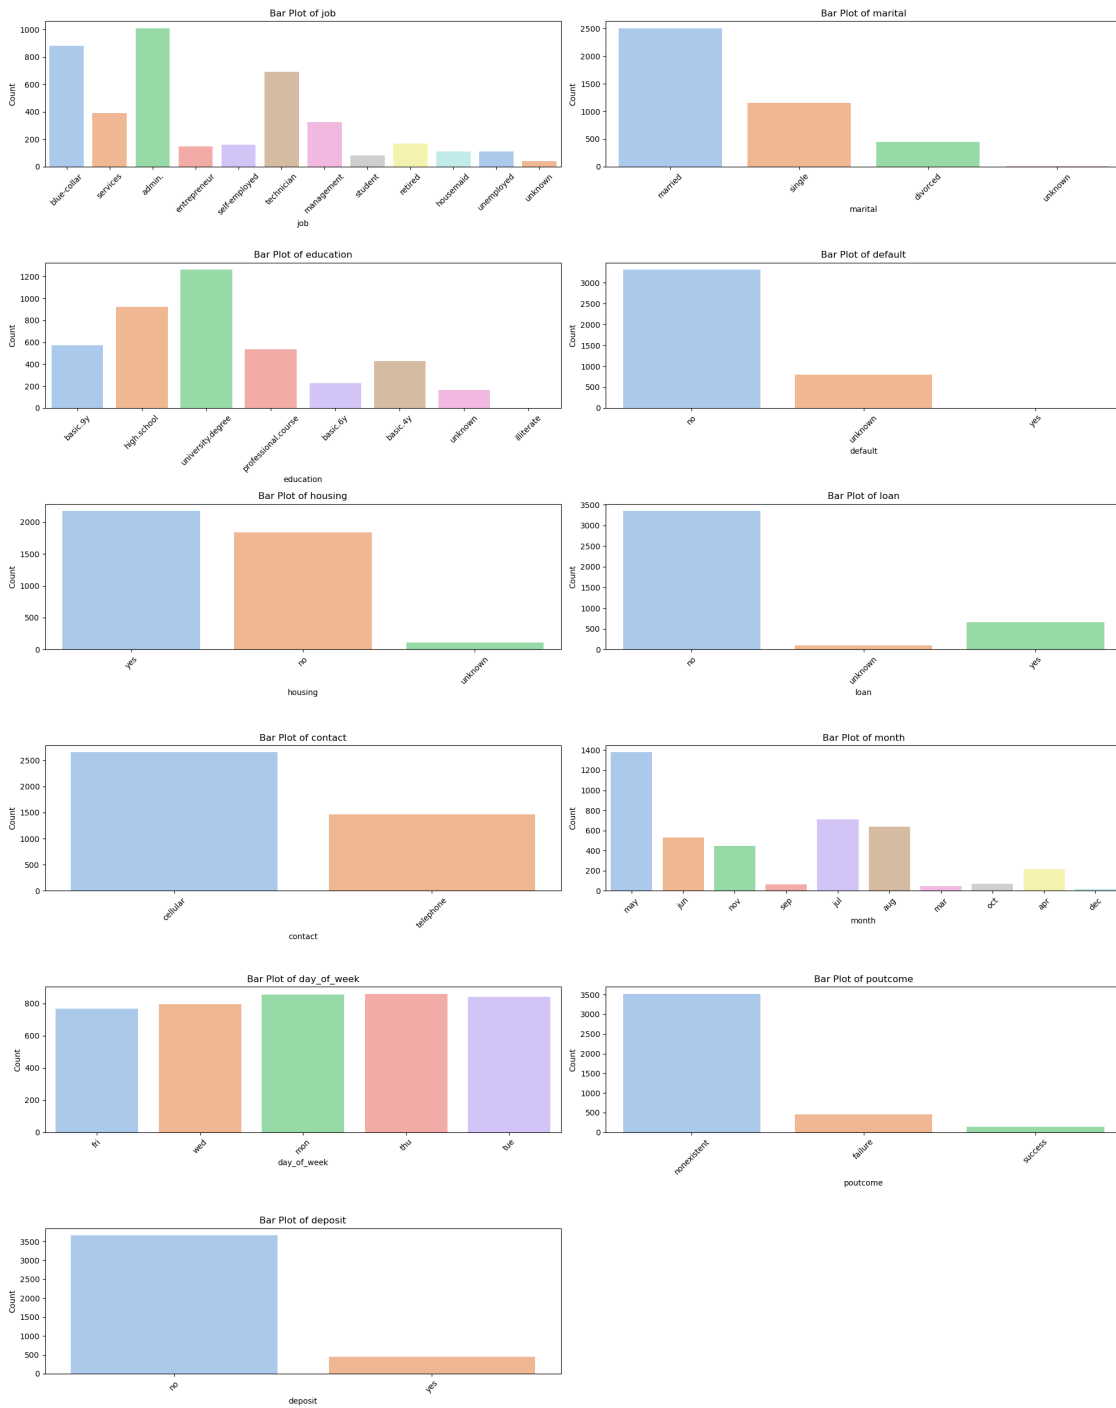


```
[65]: num_plots = len(cat_cols)
num_rows = (num_plots + 1) // 2
num_cols = 2

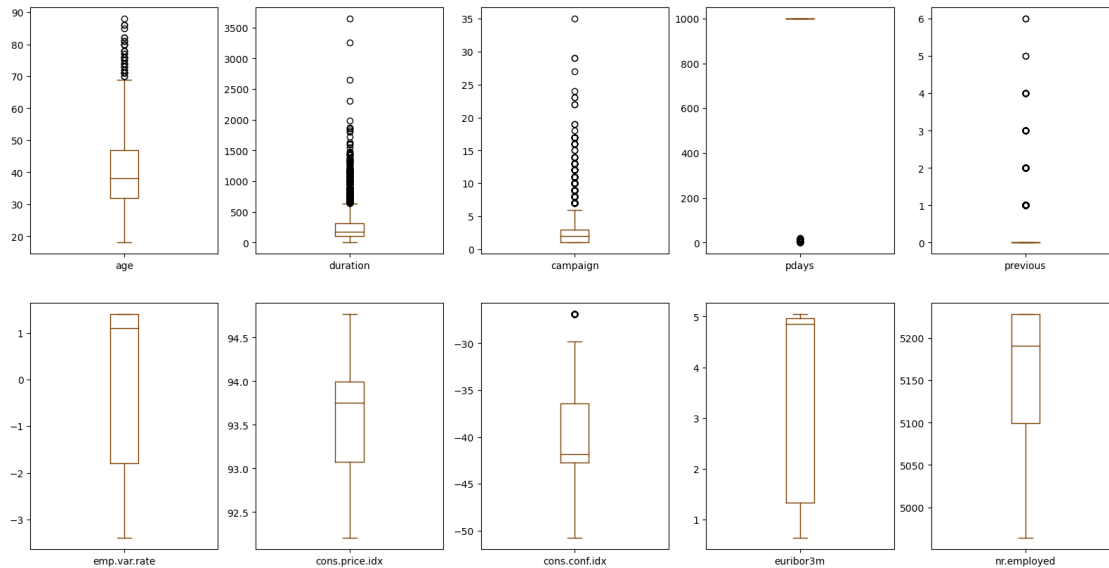
plt.figure(figsize=(20, 25))
```

```
for i, feature in enumerate(cat_cols, 1):
    plt.subplot(num_rows, num_cols, i)
    sns.countplot(x=feature, data=df, palette='pastel')
    plt.title(f'Bar Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

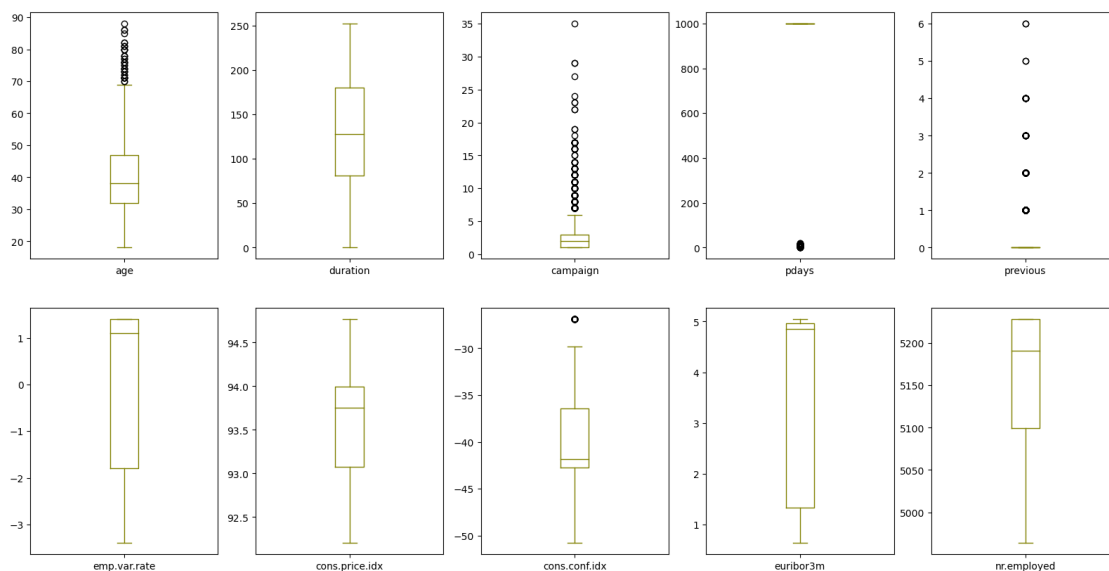


```
[33]: df.plot(kind='box', subplots=True, layout=(2,5),figsize=(20,10),color='#7b3f00')
plt.show()
```

```
[35]: column = df[['age','campaign','duration']]
q1 = np.percentile(column, 25)
q3 = np.percentile(column, 75)
iqr = q3 - q1
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr
df[['age','campaign','duration']] = column[(column > lower_bound) & (column <=
    ↪upper_bound)]
```

```
[37]: df.plot(kind='box', subplots=True, layout=(2,5),figsize=(20,10),color='#808000')
plt.show()
```




```
[39]: numeric_df = df.drop(columns=cat_cols)
corr = numeric_df.corr()
# Print the correlation matrix
print(corr)
# Filter correlations with absolute value >= 0.90
corr = corr[abs(corr) >= 0.90]
sns.heatmap(corr,annot=True,cmap='Set3',linewidths=0.2)
plt.show()
```

	age	duration	campaign	pdays	previous	\
age	1.000000	0.014048	-0.014169	-0.043425	0.050931	
duration	0.014048	1.000000	-0.218111	-0.093694	0.094206	
campaign	-0.014169	-0.218111	1.000000	0.058742	-0.091490	
pdays	-0.043425	-0.093694	0.058742	1.000000	-0.587941	
previous	0.050931	0.094206	-0.091490	-0.587941	1.000000	
emp.var.rate	-0.019192	-0.063870	0.176079	0.270684	-0.415238	
cons.price.idx	-0.000482	-0.013338	0.145021	0.058472	-0.164922	
cons.conf.idx	0.098135	0.045889	0.007882	-0.092090	-0.051420	
euribor3m	-0.015033	-0.067815	0.159435	0.301478	-0.458851	
nr.employed	-0.041936	-0.097339	0.161037	0.381983	-0.514853	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	\
age	-0.019192	-0.000482	0.098135	-0.015033	
duration	-0.063870	-0.013338	0.045889	-0.067815	
campaign	0.176079	0.145021	0.007882	0.159435	
pdays	0.270684	0.058472	-0.092090	0.301478	
previous	-0.415238	-0.164922	-0.051420	-0.458851	
emp.var.rate	1.000000	0.755155	0.195022	0.970308	
cons.price.idx	0.755155	1.000000	0.045835	0.657159	
cons.conf.idx	0.195022	0.045835	1.000000	0.276595	
euribor3m	0.970308	0.657159	0.276595	1.000000	
nr.employed	0.897173	0.472560	0.107054	0.942589	

	nr.employed
age	-0.041936
duration	-0.097339
campaign	0.161037
pdays	0.381983
previous	-0.514853
emp.var.rate	0.897173
cons.price.idx	0.472560
cons.conf.idx	0.107054
euribor3m	0.942589
nr.employed	1.000000



```
[41]: high_corr_cols = ['emp.var.rate', 'euribor3m', 'nr.employed']
```

```
[43]: df1 = df.copy()
df1.columns
```

```
[43]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
          'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
          'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
          'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
          dtype='object')
```

```
[45]: df1.drop(high_corr_cols, inplace=True, axis=1) # axis=1 indicates columns
df1.columns
```

```
[45]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
          'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
          'previous', 'poutcome', 'cons.price.idx', 'cons.conf.idx', 'deposit'],
          dtype='object')
```

```
[47]: df1.shape
```

```
[47]: (4119, 18)
```

```
[49]: from sklearn.preprocessing import LabelEncoder
lb = LabelEncoder()
df_encoded = df1.apply(lb.fit_transform)
df_encoded
```

```
[49]:
```

	age	job	marital	education	default	housing	loan	contact	month	\
0	12	1	1	2	0	2	0	0	6	
1	21	7	2	3	0	0	0	1	6	
2	7	7	1	3	0	2	0	1	4	
3	20	7	1	2	0	1	1	1	4	
4	29	0	1	6	0	2	0	0	7	
...
4114	12	0	1	1	0	2	2	0	3	
4115	21	0	1	3	0	2	0	1	3	
4116	9	8	2	3	0	0	0	0	6	
4117	40	0	1	3	0	0	0	0	1	
4118	16	4	2	3	0	2	0	0	7	

	day_of_week	duration	campaign	pdays	previous	poutcome	\
0	0	250	1	20	0	1	
1	0	250	3	20	0	1	
2	4	224	0	20	0	1	
3	0	14	2	20	0	1	
4	1	55	0	20	0	1	
...
4114	2	50	0	20	0	1	
4115	0	216	0	20	0	1	
4116	1	61	1	20	1	0	
4117	0	250	0	20	0	1	
4118	4	172	0	20	0	1	

	cons.price.idx	cons.conf.idx	deposit
0	8	4	0
1	18	16	0
2	23	8	0
3	23	8	0
4	11	7	0
...
4114	17	6	0
4115	17	6	0
4116	8	4	0
4117	13	17	0
4118	11	7	0

```
[4119 rows x 18 columns]
```

```
[51]: df_encoded['deposit'].value_counts()
```

```
[51]: deposit
0      3668
1       451
Name: count, dtype: int64
```

```
[53]: x = df_encoded.drop('deposit',axis=1)  # independent variable
y = df_encoded['deposit']                  # dependent variable
print(x.shape)
print(y.shape)
print(type(x))
print(type(y))
```

```
(4119, 17)
(4119,)
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

```
[27]: from sklearn.model_selection import train_test_split
```

```
[28]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
↳25,random_state=1)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(3089, 17)
(1030, 17)
(3089,)
(1030,)
```

```
[83]: from sklearn.metrics import
↳confusion_matrix,classification_report,accuracy_score

def eval_model(y_test,y_pred):
    acc = accuracy_score(y_test,y_pred)
    print('Accuracy_Score',acc)
    cm = confusion_matrix(y_test,y_pred)
    print('Confusion Matrix\n',cm)
    print('Classification Report\n',classification_report(y_test,y_pred))

def mscore(model):
```

```

train_score = model.score(x_train,y_train)
test_score = model.score(x_test,y_test)
print('Training Score',train_score)
print('Testing Score',test_score)

```

```
[31]: mscore(dt)
```

```

Training Score 0.9148591777274199
Testing Score 0.8990291262135922

```

```
[32]: ypred_dt = dt.predict(x_test)
print(ypred_dt)
```

```
[0 0 1 ... 0 0 0]
```

```
[33]: eval_model(y_test,ypred_dt)
```

```

Accuracy_Score 0.8990291262135922
Confusion Matrix
[[905  25]
 [ 79  21]]
Classification Report

```

	precision	recall	f1-score	support
0	0.92	0.97	0.95	930
1	0.46	0.21	0.29	100
accuracy			0.90	1030
macro avg	0.69	0.59	0.62	1030
weighted avg	0.87	0.90	0.88	1030

```
[34]: from sklearn.tree import plot_tree
```

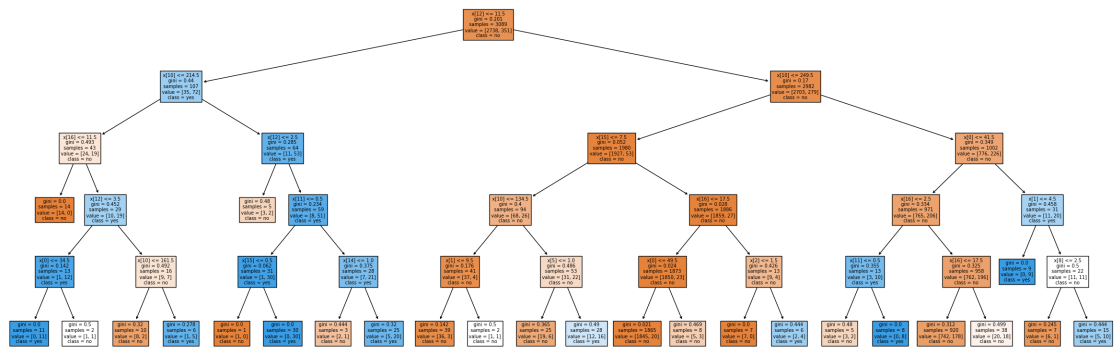
```
[35]: cn = ['no','yes']
fn = x_train.columns
print(fn)
print(cn)
```

```

Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'cons.price.idx', 'cons.conf.idx'],
      dtype='object')
['no', 'yes']

```

```
[79]: plt.figure(figsize=(30,10))
plot_tree(dt,class_names=cn,filled=True)
plt.show()
```



```
[45]: mscore(dt1)
```

Training Score 0.9080608611201036
 Testing Score 0.9048543689320389

```
[46]: ypred_dt1 = dt1.predict(x_test)
```

```
[47]: eval_model(y_test,ypred_dt1)
```

Accuracy_Score 0.9048543689320389
 Confusion Matrix
 [[915 15]
 [83 17]]

Classification Report

	precision	recall	f1-score	support
0	0.92	0.98	0.95	930
1	0.53	0.17	0.26	100
accuracy			0.90	1030
macro avg	0.72	0.58	0.60	1030
weighted avg	0.88	0.90	0.88	1030